

```

library(tree)      # tree(), cv.tree()
library(ISLR)      # OJ dataframe
library(ggplot2)

# question 1

str(OJ)

## 'data.frame':  1070 obs. of  18 variables:
## $ Purchase      : Factor w/ 2 levels "CH","MM": 1 1 1 2 1 1 1 1 1 1 ...
## $ WeekofPurchase: num  237 239 245 227 228 230 232 234 235 238 ...
## $ StoreID       : num   1 1 1 1 7 7 7 7 7 7 ...
## $ PriceCH       : num   1.75 1.75 1.86 1.69 1.69 1.69 1.69 1.75 1.75 1.75 ...
## $ PriceMM       : num   1.99 1.99 2.09 1.69 1.69 1.99 1.99 1.99 1.99 1.99 ...
## $ DiscCH        : num    0 0 0.17 0 0 0 0 0 0 0 ...
## $ DiscMM        : num    0 0.3 0 0 0 0 0.4 0.4 0.4 0.4 ...
## $ SpecialCH     : num    0 0 0 0 0 0 1 1 0 0 ...
## $ SpecialMM     : num    0 1 0 0 0 1 1 0 0 0 ...
## $ LoyalCH       : num    0.5 0.6 0.68 0.4 0.957 ...
## $ SalePriceMM   : num    1.99 1.69 2.09 1.69 1.69 1.99 1.59 1.59 1.59 1.59 ...
## $ SalePriceCH   : num    1.75 1.75 1.69 1.69 1.69 1.69 1.69 1.75 1.75 1.75 ...
## $ PriceDiff     : num    0.24 -0.06 0.4 0 0 0.3 -0.1 -0.16 -0.16 -0.16 ...
## $ Store7        : Factor w/ 2 levels "No","Yes": 1 1 1 1 2 2 2 2 2 2 ...
## $ PctDiscMM     : num    0 0.151 0 0 0 ...
## $ PctDiscCH     : num    0 0 0.0914 0 0 ...
## $ ListPriceDiff : num    0.24 0.24 0.23 0 0 0.3 0.3 0.24 0.24 0.24 ...
## $ STORE         : num    1 1 1 1 0 0 0 0 0 0 ...

# factors
OJ$StoreID = factor(OJ$StoreID)
OJ$STORE = factor(OJ$STORE)
OJ$SpecialCH = factor(OJ$SpecialCH)
OJ$SpecialMM = factor(OJ$SpecialMM)

# training/test sets
n = nrow(OJ)
n

## [1] 1070

RNGkind(sample.kind='Rounding')

## Warning in RNGkind(sample.kind = "Rounding"): non-uniform 'Rounding' sampler
## used

set.seed(1,sample.kind = 'Rounding')

## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding' sampler
## used

train = sample(1:n, 800)
test = (1:n)[-train]      # 270 obs

# a) single classification tree
tree.OJ = tree(Purchase ~ ., data=OJ[train,] )
summary(tree.OJ)

```

```

##
## Classification tree:
## tree(formula = Purchase ~ ., data = OJ[train, ])
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SpecialCH"    "ListPriceDiff"
## Number of terminal nodes: 8
## Residual mean deviance: 0.7305 = 578.6 / 792
## Misclassification error rate: 0.165 = 132 / 800

#
# training error rate 0.165
# most important classifiers: LoyalCH, PriceDiff

# test error rate
y_hat = predict(tree.OJ,newdata=OJ[test,],type="class" ) # gives classification labels
table(y_hat,OJ[test,]$Purchase)

##
## y_hat  CH  MM
##      CH 147 49
##      MM 12 62

aux=prop.table(table( y_hat, OJ[test,]$Purchase ))
round(aux,3)

##
## y_hat      CH      MM
##      CH 0.544 0.181
##      MM 0.044 0.230

#
# test error rate 0.2259

# b) cross-validation to find optimal tree size
set.seed(2,sample.kind = 'Rounding')

## Warning in set.seed(2, sample.kind = "Rounding"): non-uniform 'Rounding' sampler
## used



cv.OJ = cv.tree(tree.OJ,FUN=prune.misclass) # CV misclass based
cv.OJ$size      #[1] 8 5 2 1

## [1] 8 5 2 1
cv.OJ$dev        #[1] 151 151 150 306

## [1] 151 151 150 306

#
par(mfrow=c(1,2))
plot(cv.OJ$dev~cv.OJ$size,type='l')
plot(cv.OJ$dev~cv.OJ$k,type='l')

```



```
par(mfrow=c(1,1))
#
# best tree with 2 terminal nodes
#
# prune training tree to 2 nodes
prune2 = prune.misclass(tree.OJ, best=2)
summary(prune2)

##
## Classification tree:
## snip.tree(tree = tree.OJ, nodes = 3:2)
## Variables actually used in tree construction:
## [1] "LoyalCH"
## Number of terminal nodes: 2
## Residual mean deviance: 0.9115 = 727.4 / 798
## Misclassification error rate: 0.1825 = 146 / 800

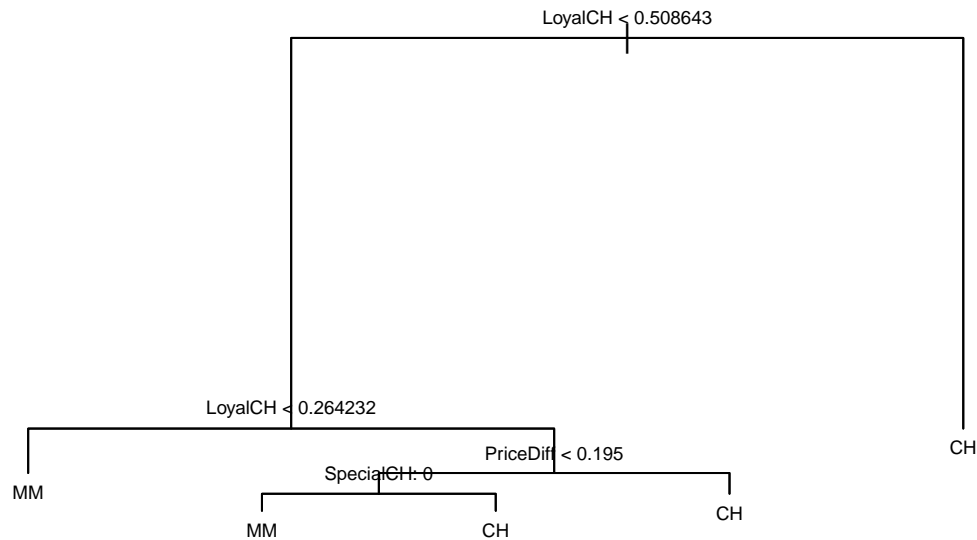
#
# LoyalCH is best classifier
#
# c) prune tree to 5 nodes

prune5b = prune.misclass(tree.OJ, best=5)
summary(prune5b)

##
## Classification tree:
## snip.tree(tree = tree.OJ, nodes = 3:4)
## Variables actually used in tree construction:
## [1] "LoyalCH" "PriceDiff" "SpecialCH"
## Number of terminal nodes: 5
## Residual mean deviance: 0.8256 = 656.4 / 795
## Misclassification error rate: 0.165 = 132 / 800
```

```
plot(prune5b)
text(prune5b,cex=0.6,pretty=0)
title("CVd tree with 5 terminal nodes")
```

## CVd tree with 5 terminal nodes



```
# test error rate
y_hat = predict( prune5b, newdata=OJ[test,], type="class" )
table( y_hat, OJ[test,]$Purchase )
```

```
##
## y_hat  CH  MM
##      CH 147 49
##      MM 12 62
```

```
#
aux=prop.table(table( y_hat, OJ[test,]$Purchase ))
round(aux,3)
```

```
##
## y_hat    CH    MM
##      CH 0.544 0.181
##      MM 0.044 0.230
```

```
#
# test error rate [1] 0.226
```

```
# d) Random Forest
library(randomForest)
set.seed(1)
rf1 = randomForest(Purchase~.,data = OJ, subset = train, importance=T)
#
ytest = OJ[test,]$Purchase
y_hat = predict(rf1, newdata=OJ[test,], type="class" )
table( y_hat, OJ[test,]$Purchase )
```

```
##
## y_hat  CH  MM
```

```
##      CH 138 33
##      MM  21 78

aux=prop.table(table( y_hat, OJ[test,]$Purchase ))
round(aux,3)

##
## y_hat      CH      MM
##      CH 0.511 0.122
##      MM 0.078 0.289

1-sum(diag(aux))

## [1] 0.2

#
# test error rate is 0.20
#
# e) Boosted Trees

# install.packages('doParallel')
library(doParallel)
cl = makePSOCKcluster(4)
registerDoParallel(cl)

# install.packages("e1071")
# install.packages("caret")

library(e1071)
library(caret) # train()
#
#
gbmGrid = expand.grid(n.minobsinnode = 10,
                     interaction.depth = c(4,7,9),
                     n.trees = seq(1200,2000,by = 200),
                     shrinkage = c(0.001, 0.01))

gbmGrid

##      n.minobsinnode interaction.depth n.trees shrinkage
## 1              10              4      1200      0.001
## 2              10              7      1200      0.001
## 3              10              9      1200      0.001
## 4              10              4      1400      0.001
## 5              10              7      1400      0.001
## 6              10              9      1400      0.001
## 7              10              4      1600      0.001
## 8              10              7      1600      0.001
## 9              10              9      1600      0.001
## 10             10              4      1800      0.001
## 11             10              7      1800      0.001
## 12             10              9      1800      0.001
## 13             10              4      2000      0.001
## 14             10              7      2000      0.001
## 15             10              9      2000      0.001
## 16             10              4      1200      0.010
## 17             10              7      1200      0.010
```

```
## 18          10          9    1200    0.010
## 19          10          4    1400    0.010
## 20          10          7    1400    0.010
## 21          10          9    1400    0.010
## 22          10          4    1600    0.010
## 23          10          7    1600    0.010
## 24          10          9    1600    0.010
## 25          10          4    1800    0.010
## 26          10          7    1800    0.010
## 27          10          9    1800    0.010
## 28          10          4    2000    0.010
## 29          10          7    2000    0.010
## 30          10          9    2000    0.010
```

```
ytrain = OJ[train,1]
xtrain = OJ[train,2:18]
set.seed(100)
gbmTune = train(xtrain, ytrain, method = "gbm", tuneGrid = gbmGrid, verbose = F)


gbmTune
```

```
## Stochastic Gradient Boosting
##
## 800 samples
## 17 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 800, 800, 800, 800, 800, ...
## Resampling results across tuning parameters:
##
## shrinkage interaction.depth n.trees Accuracy Kappa
## 0.001      4                1200    0.8219105 0.6197506
## 0.001      4                1400    0.8214937 0.6196363
## 0.001      4                1600    0.8221785 0.6217301
## 0.001      4                1800    0.8234257 0.6248542
## 0.001      4                2000    0.8227548 0.6235989
## 0.001      7                1200    0.8168915 0.6088627
## 0.001      7                1400    0.8189407 0.6140274
## 0.001      7                1600    0.8176114 0.6120073
## 0.001      7                1800    0.8182804 0.6138473
## 0.001      7                2000    0.8178823 0.6131856
## 0.001      9                1200    0.8146422 0.6038275
## 0.001      9                1400    0.8153115 0.6064219
## 0.001      9                1600    0.8159850 0.6084310
## 0.001      9                1800    0.8165043 0.6098720
## 0.001      9                2000    0.8156910 0.6084564
## 0.010      4                1200    0.8134845 0.6050466
## 0.010      4                1400    0.8109752 0.5997210
## 0.010      4                1600    0.8090566 0.5958897
## 0.010      4                1800    0.8078435 0.5933634
## 0.010      4                2000    0.8074577 0.5923253
## 0.010      7                1200    0.8057980 0.5891099
## 0.010      7                1400    0.8044611 0.5862807
```

```
## 0.010      7      1600      0.8023049 0.5816767
## 0.010      7      1800      0.8017511 0.5801109
## 0.010      7      2000      0.7994695 0.5754616
## 0.010      9      1200      0.8034277 0.5845886
## 0.010      9      1400      0.8014290 0.5799804
## 0.010      9      1600      0.8006433 0.5784710
## 0.010      9      1800      0.7989097 0.5745151
## 0.010      9      2000      0.7974069 0.5716254
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were n.trees = 1800, interaction.depth =
## 4, shrinkage = 0.001 and n.minobsinnode = 10.
```

*# This report shows the Accuracy rates for each instance in the Grid  
# it also shows the parameter values for the highest train Accuracy rate*

*# display training Accuracy rates*  
ggplot(gbmTune)



```
ytest = OJ[-train,1]
xtest = OJ[-train,2:18]

# test error rate
y_hat = predict(gbmTune,newdata=xtest, type="raw" )
table(y_hat,ytest)
```

```
##      ytest
```

```
## y_hat  CH  MM
##      CH 141  28
##      MM  18  83

aux=prop.table(table(y_hat,ytest))
round(aux,3)

##      ytest
## y_hat    CH    MM
##      CH 0.522 0.104
##      MM 0.067 0.307

1-sum(diag(aux))

## [1] 0.1703704
# test error rate is 0.17
#
#
# question 2

# a)
#
population = 234564000
d0 = read.csv('cereal.csv')
str(d0)

## 'data.frame': 1250 obs. of 3 variables:
## $ Group : int 1 2 2 4 2 2 2 3 3 4 ...
## $ Spend : num 14.77 8.15 8 9.31 12.09 ...
## $ Breakfast: int 4 4 2 1 4 4 4 4 3 3 ...

n = nrow(d0)
table(d0$Group)

##
## 1 2 3 4
## 269 484 241 256

x = 269
prop.test(x,n)$conf.int

## [1] 0.1929252 0.2392505
## attr(,"conf.level")
## [1] 0.95

population*prop.test(x,n)$conf.int

## [1] 45253302 56119555
## attr(,"conf.level")
## [1] 0.95

#
# the number of Americans who are concerned about eating healthy foods is
# somewhere between 45253302 and 56119555.
#
# b)
head(d0)
```



```
##   Group Spend Breakfast
## 1     1 14.77          4
## 2     2  8.15          4
## 3     2  8.00          2
## 4     4  9.31          1
## 5     2 12.09          4
## 6     2  7.13          4
```

```
d0$Group[d0$Group != 1] = 0
head(d0)
```

```
##   Group Spend Breakfast
## 1     1 14.77          4
## 2     0  8.15          4
## 3     0  8.00          2
## 4     0  9.31          1
## 5     0 12.09          4
## 6     0  7.13          4
```

```
table(d0$Group)
```

```
##
##    0    1
## 981 269
```

```
#
# Average Spending by Group
tapply(d0$Spend,d0$Group,mean)
```

```
##          0          1
## 10.59541 11.57926
```

```
# test Ho: average Spending Group 1 = average Spending Group 0 (all others)
Group1 = d0$Spend[d0$Group == 1]
Group0 = d0$Spend[d0$Group == 0]
t.test(Group1,Group0,'greater')
```

```
##
## Welch Two Sample t-test
##
## data:  Group1 and Group0
## t = 4.146, df = 643.95, p-value = 1.918e-05
## alternative hypothesis: true difference in means is greater than 0
## 95 percent confidence interval:
##  0.5929601      Inf
## sample estimates:
## mean of x mean of y
## 11.57926 10.59541
```

```
#
# based on p-value reject Ho in favor of Ha: true difference in means is greater than 0
# conclude that Americans concerned about eating healthy foods spend on average, more
# than other Americans
#
#
#
```

```
# question 3
```

```
#  
## ID Prediction  
## 1 1156.0714  
## 2 141.7590  
## 3 709.4666  
## 4 746.4439  
## 5 141.7590  
## 6 141.7590  
## 7 510.0157  
## 8 141.7590  
## 9 510.0157
```