

ISE 599 Midterm

Shaoqian Chen 8831737894 Wednesday

3/4/2020

1. The OJ data set from the ISLR library contains 1070 purchases where the customer either purchased

Citrus Hill or Minute Maid Orange Juice (Use ?OJ for more details). It is of interest to predict

Purchase using all other variables. Use `set.seed(1)` to create a training set containing a random

sample of 800 observations, and a test set containing the remaining observations.

Fit a classification tree to the training data to answer questions (a) to (c).

a

Plot the tree. What is the training error rate? What is the test error rate?

```
library(ISLR)
library(tree)
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
library(gbm)
```

```
## Loaded gbm 2.1.5
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
##  
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:randomForest':  
##  
##      margin
```

```
str(OJ)
```

```
## 'data.frame':    1070 obs. of  18 variables:  
## $ Purchase      : Factor w/ 2 levels "CH","MM": 1 1 1 2 1 1 1 1 1 1 ...  
## $ WeekofPurchase: num  237 239 245 227 228 230 232 234 235 238 ...  
## $ StoreID       : num   1 1 1 1 7 7 7 7 7 7 ...  
## $ PriceCH       : num   1.75 1.75 1.86 1.69 1.69 1.69 1.69 1.75 1.75 1.75 ...  
## $ PriceMM       : num   1.99 1.99 2.09 1.69 1.69 1.99 1.99 1.99 1.99 1.99 ...  
## $ DiscCH        : num   0 0 0.17 0 0 0 0 0 0 0 ...  
## $ DiscMM        : num   0 0.3 0 0 0 0 0.4 0.4 0.4 0.4 ...  
## $ SpecialCH     : num   0 0 0 0 0 0 1 1 0 0 ...  
## $ SpecialMM     : num   0 1 0 0 0 1 1 0 0 0 ...  
## $ LoyalCH       : num   0.5 0.6 0.68 0.4 0.957 ...  
## $ SalePriceMM   : num   1.99 1.69 2.09 1.69 1.69 1.99 1.59 1.59 1.59 1.59 ...  
## $ SalePriceCH   : num   1.75 1.75 1.69 1.69 1.69 1.69 1.69 1.75 1.75 1.75 ...  
## $ PriceDiff     : num   0.24 -0.06 0.4 0 0 0.3 -0.1 -0.16 -0.16 -0.16 ...  
## $ Store7        : Factor w/ 2 levels "No","Yes": 1 1 1 1 2 2 2 2 2 2 ...  
## $ PctDiscMM     : num   0 0.151 0 0 0 ...  
## $ PctDiscCH     : num   0 0 0.0914 0 0 ...  
## $ ListPriceDiff : num   0.24 0.24 0.23 0 0 0.3 0.3 0.24 0.24 0.24 ...  
## $ STORE         : num   1 1 1 1 0 0 0 0 0 0 ...
```

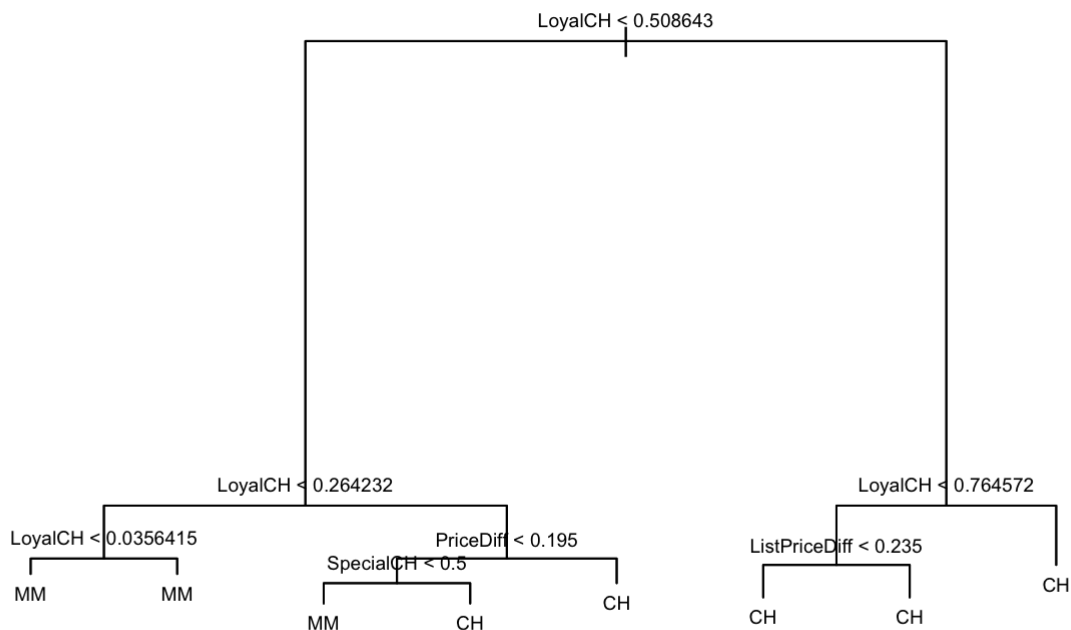
```
#RNGkind(sample.kind = 'Rounding')  
set.seed(1)  
train = sample(nrow(OJ),800)  
OJ.train = OJ[train, ]  
OJ.test = OJ[-train, ]
```

```
tree.OJ = tree(Purchase ~ . , data = OJ, subset = train)  
summary(tree.OJ)
```

```
##
## Classification tree:
## tree(formula = Purchase ~ ., data = OJ, subset = train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SpecialCH"    "ListPriceDiff"
## Number of terminal nodes: 8
## Residual mean deviance: 0.7305 = 578.6 / 792
## Misclassification error rate: 0.165 = 132 / 800
```

```
plot(tree.OJ)
text(tree.OJ,cex=0.6,pretty=0)
title("Tree from the training set")
```

Tree from the training set



```
#test error rate
OJ.pred= predict(tree.OJ, OJ.test, type = "class")
table(OJ.test$Purchase, OJ.pred)
```

```
##      OJ.pred
##      CH  MM
## CH 147  12
## MM  49  62
```

```
mean(OJ.test$Purchase != OJ.pred)
```

```
## [1] 0.2259259
```

Training error rate = 16.5%;

Test error rate = 22.59259%

b

Use `set.seed(2)` and cross-validation to find the best number of terminal nodes.

Which tree size corresponds to the lowest cross-validated classification error rate?



```
set.seed(2)
tree1= cv.tree(tree.OJ, FUN = prune.misclass)
names(tree1)
```

```
## [1] "size" "dev" "k" "method"
```

```
tree1
```

```
## $size
## [1] 8 5 2 1
##
## $dev
## [1] 151 151 150 306
##
## $k
## [1] -Inf 0.000000 4.666667 160.000000
##
## $method
## [1] "misclass"
##
## attr(,"class")
## [1] "prune" "tree.sequence"
```

```
par(mfrow=c(1,2))
plot(tree1$size, tree1$dev, type = "b", xlab = "Tree Size", ylab = "Deviance")
plot(tree1$dev, tree1$k, type = "b")
```



####

Tree size with 2 nodes corresponding to the lowest cross-validated classification error rate


C

Plot a pruned tree with five terminal nodes. What is the test error rate?

```
pruned1= prune.tree(tree.OJ, best = 5)
summary(pruned1)
```

```
##
## Classification tree:
## snip.tree(tree = tree.OJ, nodes = 4:5)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "ListPriceDiff"
## Number of terminal nodes: 5
## Residual mean deviance: 0.7829 = 622.4 / 795
## Misclassification error rate: 0.1825 = 146 / 800
```

```
plot(pruned1)
text(pruned1)
```



```
pred2= predict(pruned1, OJ.test, type = "class")
table(OJ.test$Purchase, pred2)
```

```
##      pred2
##      CH  MM
## CH 119  40
## MM   30  81
```

```
mean(OJ.test$Purchase != pred2)
```

```
## [1] 0.2592593
```

Test error rate is 25.92593%

d) (10 pts.) Which predictors are the most important? What is the test error rate? When fitting a boosted tree the number of trees, depth of trees, shrinkage, should be carefully selected. If a tuning grid is defined, the train function can be used to tune these parameters (see p217, handout).

```
train2=randomForest(Purchase~.,data=OJ,subset=train,mtry=17,importance = T)
summary(train2)
```

```
##              Length Class  Mode
## call              6  -none-  call
## type              1  -none- character
## predicted         800  factor numeric
## err.rate         1500  -none- numeric
## confusion          6  -none- numeric
## votes            1600 matrix numeric
## oob.times         800  -none- numeric
## classes           2  -none- character
## importance         68  -none- numeric
## importanceSD       51  -none- numeric
## localImportance    0  -none-  NULL
## proximity          0  -none-  NULL
## ntree             1  -none- numeric
## mtry              1  -none- numeric
## forest            14  -none-  list
## y                 800  factor numeric
## test              0  -none-  NULL
## inbag             0  -none-  NULL
## terms             3   terms   call
```

```
pred3= predict(train2, OJ.test,type = "class")
table(OJ.test$Purchase, pred3)
```

```
##      pred3
##           CH  MM
## CH 132   27
## MM  29   82
```

```
mean(OJ.test$Purchase != pred3)
```

```
## [1] 0.2074074
```

```
importance(train2)
```

	CH	MM	MeanDecreaseAccuracy	MeanDecreaseGini
## WeekofPurchase	13.5795438	4.783217	14.412380	34.224017
## StoreID	8.0761610	13.166715	15.948238	10.793909
## PriceCH	7.7672071	3.189580	8.843823	4.529538
## PriceMM	3.2181449	5.801087	6.652419	3.928933
## DiscCH	-1.2936648	6.297739	4.761091	2.094399
## DiscMM	7.2990487	2.036354	7.795884	2.840041
## SpecialCH	6.1797561	2.362523	6.413429	5.830216
## SpecialMM	-0.6595525	-3.236131	-3.034125	2.100163
## LoyalCH	78.6907996	110.331214	120.096993	233.811843
## SalePriceMM	3.4815395	7.505012	8.435593	8.753717
## SalePriceCH	8.5697938	2.381339	8.517221	5.741133
## PriceDiff	13.6688973	14.401945	21.685324	21.911899
## Store7	0.7390921	4.155137	3.973069	1.522719
## PctDiscMM	7.1333996	2.550171	7.903262	3.064694
## PctDiscCH	0.2702835	6.305326	5.391850	2.457964
## ListPriceDiff	17.1769251	3.236658	16.899851	17.608657
## STORE	6.7506902	16.343647	17.864219	9.497779

LoyalCH is the most important predictor, test error rate is 20.74074%.

e) (10 pts.) Fit a boosted tree selecting the best parameter values. What is the test error rate?

```
gbmGrid = expand.grid(.n.minobsinnode = 10, .interaction.depth = seq(1,7,by = 2), .n.trees = seq(100,1000,by = 50),.shrinkage = c(0.01,0.1))
set.seed(100)
ytrain = OJ.train$Purchase
gbmTune = train(OJ.train[,-1],ytrain,method = 'gbm',tuneGrid = gbmGrid,verbose = FALSE)
```

```
gbmTune
```



```

## Stochastic Gradient Boosting
##
## 800 samples
## 17 predictor
## 2 classes: 'CH', 'MM'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 800, 800, 800, 800, 800, 800, ...
## Resampling results across tuning parameters:
##
##  shrinkage  interaction.depth  n.trees  Accuracy  Kappa
##  0.01      1                100      0.8126448  0.6023160
##  0.01      1                150      0.8152637  0.6095701
##  0.01      1                200      0.8167365  0.6127686
##  0.01      1                250      0.8193767  0.6182519
##  0.01      1                300      0.8203167  0.6198000
##  0.01      1                350      0.8217974  0.6225704
##  0.01      1                400      0.8223281  0.6236185
##  0.01      1                450      0.8231352  0.6249320
##  0.01      1                500      0.8238163  0.6262516
##  0.01      1                550      0.8237836  0.6257345
##  0.01      1                600      0.8229956  0.6237809
##  0.01      1                650      0.8225685  0.6230099
##  0.01      1                700      0.8224265  0.6225570
##  0.01      1                750      0.8224336  0.6225724
##  0.01      1                800      0.8224355  0.6226222
##  0.01      1                850      0.8224490  0.6225704
##  0.01      1                900      0.8223038  0.6222319
##  0.01      1                950      0.8223075  0.6223010
##  0.01      1               1000      0.8215190  0.6208595
##  0.01      3                100      0.8201629  0.6159922
##  0.01      3                150      0.8208651  0.6189627
##  0.01      3                200      0.8216441  0.6211115
##  0.01      3                250      0.8226003  0.6234273
##  0.01      3                300      0.8237094  0.6258893
##  0.01      3                350      0.8236863  0.6258643
##  0.01      3                400      0.8251697  0.6293162
##  0.01      3                450      0.8247903  0.6283911
##  0.01      3                500      0.8234409  0.6255892
##  0.01      3                550      0.8233144  0.6252757
##  0.01      3                600      0.8210015  0.6205556
##  0.01      3                650      0.8209976  0.6208720
##  0.01      3                700      0.8208282  0.6205073
##  0.01      3                750      0.8194877  0.6178076
##  0.01      3                800      0.8184110  0.6155573
##  0.01      3                850      0.8177266  0.6141966
##  0.01      3                900      0.8176231  0.6139424
##  0.01      3                950      0.8174567  0.6135368
##  0.01      3               1000      0.8166941  0.6122986
##  0.01      5                100      0.8197373  0.6138201
##  0.01      5                150      0.8208315  0.6183149
##  0.01      5                200      0.8209848  0.6196734

```

##	0.01	5	250	0.8204666	0.6190242
##	0.01	5	300	0.8196539	0.6172704
##	0.01	5	350	0.8199135	0.6179549
##	0.01	5	400	0.8200224	0.6182694
##	0.01	5	450	0.8196449	0.6177171
##	0.01	5	500	0.8177505	0.6137580
##	0.01	5	550	0.8181572	0.6148982
##	0.01	5	600	0.8178796	0.6144737
##	0.01	5	650	0.8174796	0.6134921
##	0.01	5	700	0.8165297	0.6113271
##	0.01	5	750	0.8153202	0.6088914
##	0.01	5	800	0.8152297	0.6089080
##	0.01	5	850	0.8140269	0.6064283
##	0.01	5	900	0.8137724	0.6058923
##	0.01	5	950	0.8130690	0.6044259
##	0.01	5	1000	0.8122638	0.6026553
##	0.01	7	100	0.8150343	0.6031882
##	0.01	7	150	0.8161335	0.6081195
##	0.01	7	200	0.8173236	0.6118925
##	0.01	7	250	0.8154248	0.6084090
##	0.01	7	300	0.8158179	0.6096941
##	0.01	7	350	0.8148888	0.6078400
##	0.01	7	400	0.8135864	0.6051312
##	0.01	7	450	0.8136979	0.6058094
##	0.01	7	500	0.8133259	0.6050940
##	0.01	7	550	0.8130395	0.6046176
##	0.01	7	600	0.8125303	0.6037894
##	0.01	7	650	0.8122349	0.6034510
##	0.01	7	700	0.8112841	0.6014884
##	0.01	7	750	0.8099307	0.5984575
##	0.01	7	800	0.8082035	0.5948043
##	0.01	7	850	0.8088841	0.5961994
##	0.01	7	900	0.8090183	0.5962857
##	0.01	7	950	0.8095526	0.5975911
##	0.01	7	1000	0.8075258	0.5933824
##	0.10	1	100	0.8209658	0.6195814
##	0.10	1	150	0.8240673	0.6267431
##	0.10	1	200	0.8195954	0.6176686
##	0.10	1	250	0.8178754	0.6138461
##	0.10	1	300	0.8174594	0.6133033
##	0.10	1	350	0.8156283	0.6096862
##	0.10	1	400	0.8177487	0.6141740
##	0.10	1	450	0.8151349	0.6084526
##	0.10	1	500	0.8147162	0.6081483
##	0.10	1	550	0.8131441	0.6046098
##	0.10	1	600	0.8134169	0.6047599
##	0.10	1	650	0.8122277	0.6025311
##	0.10	1	700	0.8112592	0.5998015
##	0.10	1	750	0.8113880	0.6003380
##	0.10	1	800	0.8114174	0.6003453
##	0.10	1	850	0.8087160	0.5943290
##	0.10	1	900	0.8071255	0.5913198
##	0.10	1	950	0.8070055	0.5907607
##	0.10	1	1000	0.8098409	0.5964215

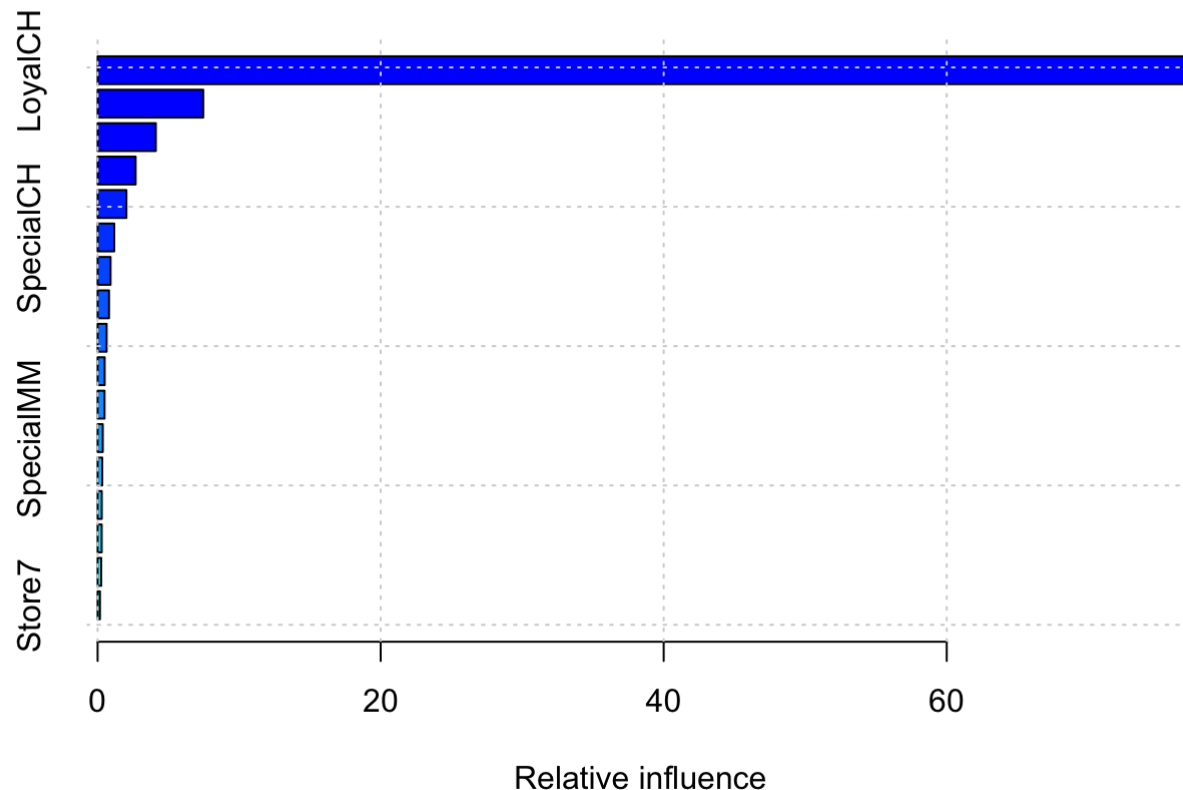
##	0.10	3	100	0.8191453	0.6171008
##	0.10	3	150	0.8136199	0.6054119
##	0.10	3	200	0.8098641	0.5972038
##	0.10	3	250	0.8073145	0.5916498
##	0.10	3	300	0.8047638	0.5861040
##	0.10	3	350	0.8041805	0.5844813
##	0.10	3	400	0.8001504	0.5762865
##	0.10	3	450	0.8008055	0.5773875
##	0.10	3	500	0.8012323	0.5786874
##	0.10	3	550	0.7999920	0.5761827
##	0.10	3	600	0.8000481	0.5763654
##	0.10	3	650	0.7980568	0.5721807
##	0.10	3	700	0.7959195	0.5675869
##	0.10	3	750	0.7972588	0.5704880
##	0.10	3	800	0.7949569	0.5659828
##	0.10	3	850	0.7950661	0.5659688
##	0.10	3	900	0.7969301	0.5702408
##	0.10	3	950	0.7957497	0.5676135
##	0.10	3	1000	0.7954923	0.5671848
##	0.10	5	100	0.8122057	0.6028173
##	0.10	5	150	0.8083589	0.5946402
##	0.10	5	200	0.8042848	0.5859704
##	0.10	5	250	0.8044721	0.5862229
##	0.10	5	300	0.8013718	0.5797095
##	0.10	5	350	0.8009573	0.5789649
##	0.10	5	400	0.7986622	0.5738113
##	0.10	5	450	0.7983941	0.5729467
##	0.10	5	500	0.7956921	0.5671547
##	0.10	5	550	0.7955383	0.5670518
##	0.10	5	600	0.7942250	0.5643898
##	0.10	5	650	0.7951858	0.5663763
##	0.10	5	700	0.7935666	0.5631283
##	0.10	5	750	0.7922183	0.5602626
##	0.10	5	800	0.7924911	0.5614454
##	0.10	5	850	0.7901876	0.5559559
##	0.10	5	900	0.7911353	0.5580953
##	0.10	5	950	0.7922358	0.5603577
##	0.10	5	1000	0.7883590	0.5524622
##	0.10	7	100	0.8056272	0.5882130
##	0.10	7	150	0.8038788	0.5843220
##	0.10	7	200	0.8008446	0.5778369
##	0.10	7	250	0.7990711	0.5744364
##	0.10	7	300	0.7973362	0.5706409
##	0.10	7	350	0.7961146	0.5678503
##	0.10	7	400	0.7970469	0.5699733
##	0.10	7	450	0.7967633	0.5692359
##	0.10	7	500	0.7956256	0.5668483
##	0.10	7	550	0.7953777	0.5662241
##	0.10	7	600	0.7943265	0.5642705
##	0.10	7	650	0.7952270	0.5661568
##	0.10	7	700	0.7952089	0.5664765
##	0.10	7	750	0.7949540	0.5660101
##	0.10	7	800	0.7941144	0.5639014
##	0.10	7	850	0.7942835	0.5644148

```
##      0.10      7      900      0.7940017  0.5639658
##      0.10      7      950      0.7926924  0.5608502
##      0.10      7     1000      0.7921381  0.5597924
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were n.trees = 400,
## interaction.depth = 3, shrinkage = 0.01 and n.minobsinnode = 10.
```

```
boost1 = gbm(Purchase~.,data = OJ.train,distribution = "gaussian", n.trees = 400, interaction.depth = 3, shrinkage = 0.01, n.minobsinnode = 10)
summary(boost1)
```

```
##              var      rel.inf
## LoyalCH      LoyalCH 77.4896890
## PriceDiff    PriceDiff 7.4636900
## ListPriceDiff ListPriceDiff 4.1137062
## StoreID      StoreID 2.6859843
## WeekofPurchase WeekofPurchase 2.0340355
## SpecialCH    SpecialCH 1.1741960
## SalePriceMM  SalePriceMM 0.9147030
## STORE        STORE 0.8066566
## DiscCH       DiscCH 0.6369599
## PriceMM      PriceMM 0.4977882
## SalePriceCH  SalePriceCH 0.4887161
## SpecialMM    SpecialMM 0.3653863
## PctDiscMM    PctDiscMM 0.3279045
## PriceCH      PriceCH 0.2934163
## DiscMM       DiscMM 0.2879382
## PctDiscCH    PctDiscCH 0.2543679
## Store7       Store7 0.1648620
```

```
grid()
```



```
pred5=predict(boost1,newdata=OJ[-train,],n.trees=400)
pred5test = ifelse(pred5>1.5,2,1)
table(OJ.test$Purchase,pred5test)
```

```
##      pred5test
##           1    2
##    CH 139   20
##    MM  25   86
```

```
result = table(OJ.test$Purchase,pred5test)
```

```
result5 =(result["MM",1]+result["CH",2])/(sum(result))
result5
```

```
## [1] 0.1666667
```

The final values used for the model were $n.trees = 400$, $interaction.depth = 3$, $shrinkage = 0.01$ and $n.minobsinnode = 10$.

test error rate is 16.66667%.

##2. In segmenting the market, a breakfast cereal manufacturer uses health and diet consciousness as the segmentation variable. Four segments are developed: ### 1 = Concerned about eating healthy foods ### 2 = Concerned primarily about weight ### 3 = Concerned about health because of illness ### 4 = Unconcerned ###

To distinguish between groups, a survey is conducted (see cereal.csv). In the survey, people are categorized as belonging to one of these groups. The most recent census reveals that 234,564,000 Americans are 18 and older.

a) (20 pts.) Use the prop.test function to find a 95% Confidence interval for the true proportion of American adults who are concerned about eating healthy foods. Then use it to estimate how many American adults belong to group 1.

```
data2 <- read.csv("cereal.csv")
head(data2)
```

```
##      Group Spend Breakfast
## 1         1  14.77         4
## 2         2   8.15         4
## 3         2   8.00         2
## 4         4   9.31         1
## 5         2  12.09         4
## 6         2   7.13         4
```

```
table(data2$Group)
```

```
##
##      1      2      3      4
## 269 484 241 256
```

```
dim(data2)
```

```
## [1] 1250      3
```

```
prop.test(269,1250)
```

```
##
## 1-sample proportions test with continuity correction
##
## data: 269 out of 1250, null probability 0.5
## X-squared = 404.42, df = 1, p-value < 2.2e-16
## alternative hypothesis: true p is not equal to 0.5
## 95 percent confidence interval:
##  0.1929252 0.2392505
## sample estimates:
##           p
## 0.2152
```

```
lower = 234564000 * 0.1929252
upper = 234564000 * 0.2392505
lower
```

```
## [1] 45253307
```

```
upper
```

```
## [1] 56119554
```

The 95% confidence interval of American adults in Group 1 is: [45253307,56119554]

b) (20 pts.) Each respondent was also asked the amount spent on breakfast cereal in an average month. The company would like to know whether on average the market segment concerned about eating healthy foods outspends the other market segments.

```
d2 = data2
d2[d2==3]=2
d2[d2==4]=2
```

```
aux2 = tapply(d2$Spend, d2$Group, mean)
aux2
```

```
##          1          2
## 11.57926 10.59541
```

```
abs_diff = aux2[1]-aux2[2]
abs_diff
```

```
##          1
## 0.9838437
```


```
function1 <- function(x, n1, n2)
{
  n <- n1 + n2
  idx_b <- sample(1:n, n1)
  idx_a <- setdiff(1:n, idx_b)
  mean_diff <- mean(x[idx_b]) - mean(x[idx_a])
  return(mean_diff)
}
```

```
x = d2$Spend
difference = rep(0,1000)
for(i in 1:1000) difference[i] = function1(x,269,981)
hist(difference)
abline(v=abs_diff,col = 'red')
abs_diff
```

```
##          1
## 0.9838437
```

```
text('0.9838437',x = 0.7,y = 200, col="red")
```

Histogram of difference



```
mean(difference > abs_diff)
```

```
## [1] 0
```

```
t.test(Spend~Group, data=d2, alternative = "greater")
```



```
##  
## Welch Two Sample t-test  
##  
## data: Spend by Group  
## t = 4.146, df = 643.95, p-value = 1.918e-05  
## alternative hypothesis: true difference in means is greater than 0  
## 95 percent confidence interval:  
## 0.5929601 Inf  
## sample estimates:  
## mean in group 1 mean in group 2  
## 11.57926 10.59541
```

p-values from random sampling and test on proportions agree

p-value is smaller than alpha(0.675)

Reject H_0 ;

Conclude: Group 1 is outspend all other group.

3

USC ID: 8831737894

Row 8

CHits = 42 < 450

AtBat = 185 > 147

CRBI = 9 < 114.5

Result: 141.8